


Meta-Learning with Latent Embedding Optimization

Rusu et al., ICLR 2019

Outline


- This paper model Latent Embedding Optimization (LEO), an extension over the MAML model
 - ▶ Learn a low-dimensional latent embedding of model parameters and performs optimization-based meta learning in this space
 - ▶ Provide 2 advantages over the MAML model
 - ★ initial parameters for new tasks are conditioned on the training data, which enables a task-specific starting point for adaptation
 - ★ optimizing in low-dimensional latent space is more efficient

Comparison between LEO vs MAML



- The MAML model aims to find a set of parameters that can be used for many different tasks
- The LEO model initializes a task-dependent set of parameters, updates them in a general framework, it is more desirable

LEO Algorithm and Arch



Algorithm 2 MAML for Few-Shot Supervised Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

```

1: randomly initialize  $\theta$ 
2: while not done do
3:   Sample batch of tasks  $\mathcal{T}_i \sim p(\mathcal{T})$ 
4:   for all  $\mathcal{T}_i$  do
5:     Sample  $K$  datapoints  $\mathcal{D} = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$  from  $\mathcal{T}_i$ 
6:     Evaluate  $\nabla_\theta \mathcal{L}_{\mathcal{T}_i}(f_\theta)$  using  $\mathcal{D}$  and  $\mathcal{L}_{\mathcal{T}_i}$  in Equation (2) or (3)
7:     Compute adapted parameters with gradient descent:  $\theta'_i = \theta - \alpha \nabla_\theta \mathcal{L}_{\mathcal{T}_i}(f_\theta)$ 
8:     Sample datapoints  $\mathcal{D}'_i = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$  from  $\mathcal{T}_i$  for the meta-update
9:   end for
10:  Update  $\theta \leftarrow \theta - \beta \nabla_\theta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$  using each  $\mathcal{D}'_i$  and  $\mathcal{L}_{\mathcal{T}_i}$  in Equation 2 or 3
11: end while

```

Algorithm 1 Latent Embedding Optimization

Require: Training meta-set $S^{tr} \in \mathcal{T}$

Require: Learning rates α, η

```

1: Randomly initialize  $\phi_e, \phi_r, \phi_d$ 
2: Let  $\phi = \{\phi_e, \phi_r, \phi_d, \alpha\}$ 
3: while not converged do
4:   for number of tasks in batch do
5:     Sample task instance  $\mathcal{T}_i \sim S^{tr}$ 
6:     Let  $(\mathcal{D}^{tr}, \mathcal{D}^{val}) = \mathcal{T}_i$ 
7:     Encode  $\mathcal{D}^{tr}$  to  $\mathbf{z}$  using  $g_{\phi_e}$  and  $g_{\phi_r}$ 
8:     Decode  $\mathbf{z}$  to initial params  $\theta_i$  using  $g_{\phi_d}$ 
9:     Initialize  $\mathbf{z}' = \mathbf{z}, \theta'_i = \theta_i$ 
10:    for number of adaptation steps do
11:      Compute training loss  $\mathcal{L}_{\mathcal{T}_i}^{tr}(f_{\theta'_i})$ 
12:      Perform gradient step w.r.t.  $\mathbf{z}'$ :  $\mathbf{z}' \leftarrow \mathbf{z}' - \alpha \nabla_{\mathbf{z}'} \mathcal{L}_{\mathcal{T}_i}^{tr}(f_{\theta'_i})$ 
13:      Decode  $\mathbf{z}'$  to obtain  $\theta'_i$  using  $g_{\phi_d}$ 
14:    end for
15:    Compute validation loss  $\mathcal{L}_{\mathcal{T}_i}^{val}(f_{\theta'_i})$ 
16:  end for
17:  Perform gradient step w.r.t.  $\phi$ :  $\phi \leftarrow \phi - \eta \nabla_\phi \sum_{\mathcal{T}_i} \mathcal{L}_{\mathcal{T}_i}^{val}(f_{\theta'_i})$ 
18: end while

```

LEO Model

- Encoding and Relation Network

$$\begin{aligned}\boldsymbol{\mu}_n^e, \boldsymbol{\sigma}_n^e &= \frac{1}{NK^2} \sum_{k_n=1}^K \sum_{m=1}^N \sum_{k_m=1}^K g_{\phi_r} \left(g_{\phi_e}(\mathbf{x}_n^{k_n}), g_{\phi_e}(\mathbf{x}_m^{k_m}) \right) \\ \mathbf{z}_n &\sim q(\mathbf{z}_n | \mathcal{D}_n^{tr}) = \mathcal{N}(\boldsymbol{\mu}_n^e, \text{diag}(\boldsymbol{\sigma}_n^{e^2}))\end{aligned}$$

- ▶ Encoder all examples into intermediate codes, concatenated pair-wise
- ▶ Use the relation network to learn specific code for each class, forming a Gaussian distribution
- ▶ The hidden code \mathbf{z} is drawn from the distribution

- Decoding Network

$$\begin{aligned}\boldsymbol{\mu}_n^d, \boldsymbol{\sigma}_n^d &= g_{\phi_d}(\mathbf{z}_n) \\ \mathbf{w}_n &\sim p(\mathbf{w} | \mathbf{z}_n) = \mathcal{N}(\boldsymbol{\mu}_n^d, \text{diag}(\boldsymbol{\sigma}_n^{d^2}))\end{aligned}$$

- ▶ Decode the hidden codes into distribution's parameters, forming a Gaussian distribution
- ▶ Task-specific parameters are drawn from the resulted distribution

LEO Model

- Inner-loop objective

$$\mathcal{L}_{\mathcal{T}_i}^{tr}(f_{\theta_i}) = \sum_{(\mathbf{x}, y) \in \mathcal{D}^{tr}} \left[-\mathbf{w}_y \cdot \mathbf{x} + \log \left(\sum_{j=1}^N e^{\mathbf{w}_j \cdot \mathbf{x}} \right) \right]$$

- Outer-loop objective

$$\min_{\phi_e, \phi_r, \phi_d} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \left[\mathcal{L}_{\mathcal{T}_i}^{val}(f_{\theta'_i}) + \beta D_{KL}(q(\mathbf{z}_n | \mathcal{D}_n^{tr}) || p(\mathbf{z}_n)) + \gamma ||\text{stopgrad}(\mathbf{z}'_n) - \mathbf{z}_n||_2^2 \right] + R$$

Thank you !