

Adaptive Subspaces for Few-Shot Learning

Christian Simon, Piotr Koniusz,
Richard Nock, Mehrtash Harandi
CVPR 2020

Motivation

- Formulate FSL as generating dynamic classifier
 - Learning a universal feature extractor
 - Learning to generate a classifier dynamically from limited data.
- Method
 - Generate subspace from limited data
 - Singularity Vector Decomposition
 -

FSL review

$$p(c|\mathbf{q}) = \frac{\exp(\mathbf{W}_c^\top f_\Theta(\mathbf{q}))}{\sum_{c'} \exp(\mathbf{W}_{c'}^\top f_\Theta(\mathbf{q}))} = \frac{\exp(d_c(\mathbf{q}))}{\sum_{c'} \exp(d_{c'}(\mathbf{q}))},$$




Figure 2: Various classifiers for few-shot classification. (a) Matching networks create pairwise classifiers. (b) Prototypical networks create mean classifiers based on the sample in the same class. (c) Relation networks produce non-linear classifiers. (d) Our proposed method creates classifiers using subspaces.

Subspace for few-shot

A subspace Z is represented by a basis. This paper try to generate the basis

$$\mathbb{R}^{D \times n} \ni \mathbf{B}_i = [\mathbf{b}_1, \dots, \mathbf{b}_n]; n \leq D$$

A basis for a class can be compute by matrix decomposition (e.g.SVD)

$$\mathbf{B}_i^\top \mathbf{B}_i = \mathbf{I}_n$$

Generate subspace

Sample representation

$$\tilde{\mathbf{X}}_c = [f_{\Theta}(\mathbf{x}_{c,1}) - \boldsymbol{\mu}_c, \dots, f_{\Theta}(\mathbf{x}_{c,K}) - \boldsymbol{\mu}_c]$$

where $\boldsymbol{\mu}_c = \frac{1}{K} \sum_{\mathbf{x}_i \in \mathbf{X}_c} f_{\Theta}(\mathbf{x}_i)$.

Class representation as a subspace **Pc** from **Xc** by **SVD**

Subspace classifier

Given a query instance \mathbf{q} , distance to subspace is

$$d_c(\mathbf{q}) = -\|(\mathbf{I} - \mathbf{M}_c)(f_\Theta(\mathbf{q}) - \boldsymbol{\mu}_c)\|^2$$

Distance distribution

$$p_{c,q} = p(c|\mathbf{q}) = \frac{\exp(d_c(\mathbf{q}))}{\sum_{c'} \exp(d_{c'}(\mathbf{q}))}$$

Discriminative Deep Subspace Network

Maximize subspace distance using Grassmannian geometry

$$\delta_p^2(\mathbf{P}_i, \mathbf{P}_j) = \left\| \mathbf{P}_i \mathbf{P}_i^\top - \mathbf{P}_j \mathbf{P}_j^\top \right\|_F^2 = 2n - 2\|\mathbf{P}_i^\top \mathbf{P}_j\|_F^2.$$

Final loss function

$$-\frac{1}{NM} \sum_c \log(p_{c,q}) + \lambda \sum_{i \neq j} \|\mathbf{P}_i^\top \mathbf{P}_j\|_F^2.$$

Mean refinement

Make use of unsupervised data

$$\tilde{\mu}_c = \frac{K\mu_c + \sum_i m_i f_{\Theta}(\mathbf{r}_i)}{K + \sum_i m_i}$$

Where

$$m_i = \frac{\exp(-\|f_{\Theta}(\mathbf{r}_i) - \mu_c\|^2)}{\sum_{c'} \exp(-\|f_{\Theta}(\mathbf{r}_i) - \mu_{c'}\|^2)}$$

Algorithm

Algorithm 1 Train Deep Subspace Networks

Input: Each episode \mathcal{T}_i with S and Q

- 1: $\Theta_0 \leftarrow$ random initialization
 - 2: **for** t in $\{\mathcal{T}_1, \dots, \mathcal{T}_{N_{\mathcal{T}}}\}$ **do**
 - 3: **for** k in $\{1, \dots, N\}$ **do**
 - 4: $\tilde{\mathbf{X}}_c \leftarrow \mathbf{S}_c$
 - 5: Calculate the average of the class
 - 6: Calculate mean refinement (MR) using Eq. 6
 - 7: Subtract $\tilde{\mathbf{X}}_c$ with an offset
 - 8: $[\mathcal{U}, \Sigma, \mathcal{V}^\top] \leftarrow \text{Decompose}(\tilde{\mathbf{X}}_c)$
 - 9: $\mathbf{P}_c \leftarrow \text{Truncate } \mathcal{U}_{1, \dots, n}$
 - 10: **for** q in Q **do**
 - 11: Compute $d_c(q)$ using Eq. 2
 - 12: **end for**
 - 13: **end for**
 - 14: Compute final loss \mathcal{L}_t using Eq. 5
 - 15: Update Θ using $\nabla \mathcal{L}_t$
 - 16: **end for**
-

Result

| Model | Backbone | 1-shot | 5-shot |
|------------------------|----------|------------------------------------|------------------------------------|
| Matching Nets [4] | Conv-4 | 43.56 ± 0.84 | 55.31 ± 0.73 |
| MAML [7] | Conv-4 | 48.70 ± 1.84 | 63.11 ± 0.92 |
| Reptile [48] | Conv-4 | 49.97 ± 0.32 | 65.99 ± 0.58 |
| R2-D2 [49] | Conv-4 | 48.70 ± 0.60 | 65.50 ± 0.60 |
| Prototypical Nets [20] | Conv-4 | 44.53 ± 0.76 | 65.77 ± 0.66 |
| Relation Nets [14] | Conv-4 | 50.44 ± 0.82 | 65.32 ± 0.70 |
| DSN | Conv-4 | 51.78 ± 0.96 | 68.99 ± 0.69 |
| DSN-MR | Conv-4 | 55.88 ± 0.90 | 70.50 ± 0.68 |